

# Project 2 Final Notebook

## Stakeholder and Business Problem Howard S Wright Construction Company

Howard S Wright is a well-known construction company that is prevalent within the King County area. Howard S Wright wants to know what **features** should they focus on in order to maximize their profit when building a new home in King County.

Some factors that we initially thought that Howard S Wright could focus on were:

- Square Footage (specifically the building footprint and the total living space)
- Location
- Numbers of beds and bathrooms
- Quality of the home

In this notebook, we will explore and analyze the data in order to provide the best recommendations to focus on for Howard S Wright to maximize their profit when building a home. (this might be a little redundant, please change if necessary -- I'll go over the notebook one more time before we submit)

## Exploratory Data Analysis

```
In [1]: # Imported appropriate libraries for use in this notebook and our multilinear regression
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
import statsmodels
from statsmodels.formula.api import ols
from sklearn.model_selection import train_test_split
from sklearn.dummy import DummyRegressor
from statsmodels.tools.eval_measures import rmse
from statsmodels.api import qqplot
from scipy import stats
from sklearn.preprocessing import OneHotEncoder
from folium.plugins import FastMarkerCluster
import folium
from sklearn.metrics import r2_score
import matplotlib.ticker as ticker
```

## Created functions to help with ease of access in our linear regression analysis

```
In [2]: def sm_metrics(model, y, X):
# import associated tools
```

```

from statsmodels.tools.eval_measures import rmse, mse, meanabs

# now generate predictions
ypred = model.predict(X)

# Print values
print('Metrics:')
# MAE
print(f"Mean Absolute Error: {meanabs(y, ypred):.3f}")
# MSE
print(f"Mean Squared Error: {mse(y, ypred):.3f}")
# RMSE
print(f"Root Mean Squared Error: {rmse(y, ypred):.3f}")
return

```

In [3]:

```

def sk_metrics(y, model):
    from sklearn.metrics import mean_squared_error, mean_absolute_error

    print("Metrics:")
    # R2
    print(f"R2: {r2_score(y, model):.3f}")
    # MAE
    print(f"Mean Absolute Error: {mean_absolute_error(y, model):.3f}")
    # MSE
    print(f"Mean Squared Error: {mean_squared_error(y, model):.3f}")
    # RMSE - just MSE but set squared=False
    print(f"Root Mean Squared Error: {mean_squared_error(y, model, squared=False):.3f}")
    return

```

In [4]:

```

def qq(x_test, x_train, y_test, y_train):
    # QQ plots are generally great tools for checking for normality.
    import statsmodels.api as sm
    from sklearn.linear_model import LinearRegression

    # Calculating residuals
    lr = LinearRegression()
    lr.fit(x_train, y_train)
    y_hat = lr.predict(x_test)

    residuals = y_test - y_hat

    sm.qqplot(residuals, line = 'r');
    return

```

In [5]:

```

# evaluates the Variance Inflation Factor of X_train variables

def vif(X_train):
    from statsmodels.stats.outliers_influence import variance_inflation_factor
    # defining an empty dataframe to capture the VIF scores
    vif = pd.DataFrame()

    # For each column, run a variance_inflation_factor against all other columns to get
    vif["VIF"] = [variance_inflation_factor(X_train.values, i) for i in range(len(X_train))]

    # Label the scores with their related columns

```

```

vif["features"] = X_train.columns

# print out the vif table and return
print(vif)
return

```

```

In [6]: def lr_model_metrics(x_test, x_train, y_test, y_train, cat_vars):
# One hot encoding cat vars
onehot = OneHotEncoder(sparse=False, handle_unknown = 'ignore')
x_train_cat = pd.DataFrame(onehot.fit_transform(x_train[cat_vars]))
x_train_cat.columns = onehot.get_feature_names(cat_vars)
x_test_cat = pd.DataFrame(onehot.transform(x_test[cat_vars]))
x_test_cat.columns = onehot.get_feature_names(cat_vars)

# Resetting indices to avoid joining conflicts and creation of NaN entries
x_train_cat.reset_index(drop=True, inplace=True)
x_test_cat.reset_index(drop=True, inplace=True)
x_train.reset_index(drop=True, inplace=True)
x_test.reset_index(drop=True, inplace=True)

# Combine dummied cat vars with non-cat
x_train_df = x_train_cat.join(x_train.drop(cat_vars, axis = 1))
x_test_df = x_test_cat.join(x_test.drop(cat_vars, axis = 1))

# Run Linear regression model for data
lr = LinearRegression()
model = lr.fit(x_train_df, y_train)

print('Train Data')
sk_metrics(y_train, model.predict(x_train_df.values))

print('Test Data')
sk_metrics(y_test, model.predict(x_test_df.values))
return x_train_df, x_test_df, model

```

We utilized the `kc_house_data` as well as demographics data `demographic_spatial_join.xls` that we acquired online from King County's census data. We utilized a GIS software in order to retrieve the specific data that we needed.

Performing some basic initial exploratory data analysis to see what data we're working with and if there are any missing data values that we need to take into consideration.

```

In [7]: housedf = pd.read_csv('data/kc_house_data.csv')

```

```

In [8]: housedf.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
3   bedrooms        21597 non-null  int64

```

```

4  bathrooms      21597 non-null float64
5  sqft_living    21597 non-null int64
6  sqft_lot       21597 non-null int64
7  floors         21597 non-null float64
8  waterfront     19221 non-null object
9  view           21534 non-null object
10 condition      21597 non-null object
11 grade          21597 non-null object
12 sqft_above     21597 non-null int64
13 sqft_basement  21597 non-null object
14 yr_built       21597 non-null int64
15 yr_renovated   17755 non-null float64
16 zipcode        21597 non-null int64
17 lat           21597 non-null float64
18 long           21597 non-null float64
19 sqft_living15  21597 non-null int64
20 sqft_lot15     21597 non-null int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```


In [9]:

```
housedf.head()
```

Out[9]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN	N
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	NO	N
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	NO	N
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	NO	N
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	NO	N

5 rows × 11 columns



In [10]:

```
housedf.columns
```

Out[10]:

```

Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
      'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
      'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
      'lat', 'long', 'sqft_living15', 'sqft_lot15'],
      dtype='object')

```

Created a few dataframes for different use cases:

- **dfinit**: used to drop columns that we were not utilizing in our model/data analysis (mostly categorical columns)
- **dfrev1**: used to drop numerical columns
- **dfrevcopy**: used for testing some modeling so that I didn't have to re-run everything when I made a mistake

Also engineered some features and added them to the original housedf:

- **sqft\_per\_bedroom**: the sqft\_living divided by bedrooms value for avg sqft per bedroom
- **footprint**: the sqft\_living divided by sqft\_lot for relative sqft of usage
- **bedbathratio**: the number of bedrooms divided by bathrooms since these values have high multicollinearity

```
In [11]: housedf.drop_duplicates(keep= 'first', inplace = True)
dfinit = housedf.drop(['id', 'date', 'yr_renovated', 'condition', 'waterfront', 'lat',
                      'sqft_above', 'sqft_basement', 'view'], axis=1).copy()
# Created duplicate dataframes so that we could work on data without accidentally affecting
dfrev1 = dfinit.drop(['sqft_living15', 'sqft_lot15'], axis=1).copy()
dfrevcopy = dfinit.drop(['sqft_living15', 'sqft_lot15'], axis=1).copy()

# Feature engineered sqft_per_bedroom as a data column since for our stakeholder, knowing
# would be useful in deciding which direction to go for homebuilding.
housedf['sqft_per_bedroom'] = housedf['sqft_living'] / housedf['bedrooms']

# Feature engineered the footprint column to visualize the relationship between sqft of
# stakeholder would know on avg what home sizes they would be working with.
housedf['footprint'] = housedf['sqft_living'] / housedf['sqft_lot']


# Feature engineered the bedbathratio since they have decently high multicollinearity -
# ratio might give us a better R^2 value
housedf['bedbathratio'] = housedf['bedrooms'] / housedf['bathrooms']

# Printed out housedf just to check that columns were added correctly
housedf
```

```
Out[11]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	NaN
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	NaN
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	NaN
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	NaN
...	...	...	...	...	...	...	...	...	...
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	NaN
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	NaN
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	NaN
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	NaN
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	NaN

21597 rows × 24 columns



Pulled kc\_zip information from the [King County Website](#). We identified the unique zipcode values and created a list with those values to use in our modeling. After creating this new dataframe, we

did initial exploratory data analysis in a similar fashion for this dataset as we did for the  
kc\_house\_data.csv

```
In [12]: kc_zips = [98178, 98125, 98028, 98136, 98074, 98053, 98003, 98198,  
                  98146, 98038, 98007, 98115, 98107, 98126, 98019, 98103,  
                  98002, 98133, 98040, 98092, 98030, 98119, 98112, 98052,  
                  98027, 98117, 98058, 98001, 98056, 98166, 98023, 98070,  
                  98148, 98105, 98042, 98008, 98059, 98122, 98144, 98004,  
                  98005, 98034, 98075, 98116, 98010, 98118, 98199, 98032,  
                  98045, 98102, 98077, 98108, 98168, 98177, 98065, 98029,  
                  98006, 98109, 98022, 98033, 98155, 98024, 98011, 98031,  
                  98106, 98072, 98188, 98014, 98055, 98039]
```

```
In [13]: df2 = pd.read_excel('data/demographic_spatial_join.xls')  
  
b_series = df2.ZCTA5CE10.isin(kc_zips)  
df_zips = df2[b_series]  
  
df_names = df_zips[['ZCTA5CE10', 'NAME']]  
df_names.head()
```

```
Out[13]:
```

	ZCTA5CE10	NAME
0	98001	Southeast
1	98001	Southwest
2	98002	Southeast
3	98003	Southeast
4	98003	Southwest

```
In [14]: df_names.NAME.value_counts()
```

```
Out[14]:
```

Northeast	33
West	29
Southeast	26
Southwest	17
Shoreline	12

Name: NAME, dtype: int64

```
In [15]: df_names
```

```
Out[15]:
```

	ZCTA5CE10	NAME
0	98001	Southeast
1	98001	Southwest
2	98002	Southeast
3	98003	Southeast
4	98003	Southwest

	ZCTA5CE10	NAME
...	...	...
138	98188	Southeast
139	98188	Southwest
141	98198	Southeast
142	98198	Southwest
143	98199	West

117 rows × 2 columns

In [16]:

```
# Dropping duplicates randomly, for the moment
df_names.drop_duplicates(subset=['ZCTA5CE10'], inplace=True)
df_names.ZCTA5CE10
```

<ipython-input-16-af7d0a158eea>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Out[16]:

```
df_names.drop_duplicates(subset=['ZCTA5CE10'], inplace=True)
0      98001
2      98002
3      98003
5      98004
6      98005
...
131    98177
135    98178
138    98188
141    98198
143    98199
Name: ZCTA5CE10, Length: 70, dtype: int64
```

In [17]:

```
# lists of zipcodes
northeast = df_names.query('NAME=="Northeast"')['ZCTA5CE10'].to_list()
west = df_names.query('NAME=="West"')['ZCTA5CE10'].to_list()
shoreline = df_names.query('NAME=="Shoreline"')['ZCTA5CE10'].to_list()
southwest = df_names.query('NAME=="Southwest"')['ZCTA5CE10'].to_list()
southeast = df_names.query('NAME=="Southeast"')['ZCTA5CE10'].to_list()
print("Northeast: ", northeast)
print("Shoreline: ", shoreline)
```

Northeast: [98004, 98005, 98006, 98007, 98008, 98011, 98014, 98019, 98024, 98027, 98028, 98029, 98033, 98034, 98038, 98039, 98040, 98045, 98052, 98053, 98056, 98059, 98065, 98072, 98074, 98075, 98077]  
Shoreline: [98125, 98133, 98155, 98177]

In [18]:

```
# Create new DF for revisions, change type
dfrev2 = dfrev1.copy()
#dfrev2 = dfrev2.astype({"zipcode": str})
# use lists to make categorical bins
dfrev2['zipcode'] = dfrev2.zipcode.replace(to_replace=northeast, value = "Northeast")
```

```
dfrev2['zipcode'] = dfrev2.zipcode.replace(to_replace=west, value = "West")
dfrev2['zipcode'] = dfrev2.zipcode.replace(to_replace=shoreline, value = "Shoreline")
dfrev2['zipcode'] = dfrev2.zipcode.replace(to_replace=southwest, value = "Southwest")
dfrev2['zipcode'] = dfrev2.zipcode.replace(to_replace=southeast, value = "Southeast")

dfrev2.zipcode.value_counts()
```

```
Out[18]: Northeast    8560
West              4442
Southeast        4126
Southwest        2866
Shoreline        1603
Name: zipcode, dtype: int64
```

## Correlation and Heatmap Visualization

Prior to running the model on our datasets, we had to check the data for any correlation or multicollinearity so that we could accomodate for that in our model.

```
In [19]: housedf.corr()
```


```
Out[19]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_ab
id	1.000000	-0.016772	0.001150	0.005162	-0.012241	-0.131911	0.018608	-0.010
price	-0.016772	1.000000	0.308787	0.525906	0.701917	0.089876	0.256804	0.60!
bedrooms	0.001150	0.308787	1.000000	0.514508	0.578212	0.032471	0.177944	0.47!
bathrooms	0.005162	0.525906	0.514508	1.000000	0.755758	0.088373	0.502582	0.68!
sqft_living	-0.012241	0.701917	0.578212	0.755758	1.000000	0.173453	0.353953	0.87!
sqft_lot	-0.131911	0.089876	0.032471	0.088373	0.173453	1.000000	-0.004814	0.18!
floors	0.018608	0.256804	0.177944	0.502582	0.353953	-0.004814	1.000000	0.52!
sqft_above	-0.010799	0.605368	0.479386	0.686668	0.876448	0.184139	0.523989	1.00!
yr_built	0.021617	0.053953	0.155670	0.507173	0.318152	0.052946	0.489193	0.42!
yr_renovated	-0.012010	0.129599	0.018495	0.051050	0.055660	0.004513	0.003535	0.02!
zipcode	-0.008211	-0.053402	-0.154092	-0.204786	-0.199802	-0.129586	-0.059541	-0.26!
lat	-0.001798	0.306692	-0.009951	0.024280	0.052155	-0.085514	0.049239	-0.00!
long	0.020672	0.022036	0.132054	0.224903	0.241214	0.230227	0.125943	0.34!
sqft_living15	-0.002701	0.585241	0.393406	0.569884	0.756402	0.144763	0.280102	0.73!
sqft_lot15	-0.138557	0.082845	0.030690	0.088303	0.184342	0.718204	-0.010722	0.19!
sqft_per_bedroom	-0.011922	0.580109	-0.024076	0.517415	0.755398	0.190583	0.290654	0.67!
footprint	0.088238	0.123063	0.026798	0.287015	0.076988	-0.252601	0.556700	0.05!
bedbathratio	-0.014457	-0.249948	0.209444	-0.653476	-0.324152	-0.050648	-0.374139	-0.31!

```
In [20]:
```



```
plt.figure(figsize = (20,15))
sns.heatmap(housedf.drop(['id'], axis = 1).corr().abs(), annot=True);
```



## Model Creation and Application

Created training and test datasets for the housedf dataframe. Utilized a 80 / 20 split.

```
In [21]: y = housedf['price']
X = housedf[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'yr_built', 'zip
            'sqft_per_bedroom', 'footprint', 'bedbathratio']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=5)
```

## Dummy Regressor

Created a Dummy Regressor on the mean for use as our baseline model.

- We used the predict function on our baseline model to check if the printed out array was filled with the same values or not.
- We then checked the score of the baseline model. We know that the score should give us a 0 since it is only accounting for the intercepts. Since this is our baseline model, we will be comparing all of our other models to this as a measurement of success.

In [22]:

```
#Dummy Regressor

basemodel = DummyRegressor(strategy = 'mean')
basemodel.fit(X_train, y_train)
basemodel.predict(X_train)
```

Out[22]:

```
array([540946.37500724, 540946.37500724, 540946.37500724, ...,
       540946.37500724, 540946.37500724, 540946.37500724])
```

In [23]:

```
basemodel.score(X_train, y_train)
```

Out[23]:

```
0.0
```

We defined the `sk_metrics` function at the beginning of our notebook, and by calling it on our `y_train` and `basemodel` we can evaluate the initial  $R^2$  and RMSE value of our `basemodel`. These are the baseline scores that we are going to be comparing our future models to in order to gauge success.

In [24]:

```
sk_metrics(y_train, basemodel.predict(X_train))
```

```
Metrics:
R2: 0.000
Mean Absolute Error: 233710.087
Mean Squared Error: 135491967490.227
Root Mean Squared Error: 368092.336
```

In [25]:

```
sk_metrics(y_test, basemodel.predict(X_test))
```

```
Metrics:
R2: -0.000
Mean Absolute Error: 236167.828
Mean Squared Error: 132800122046.262
Root Mean Squared Error: 364417.511
```

## StandardScaler and OneHotEncoder

Created a `StandardScaler` for our dataset because we had data values such as 1 bedroom vs. thousands of sqft. We wanted to see a more accurate representation of our dataset as a whole.

In [26]:

```
X_train_plus_y = pd.concat((X_train, y_train), axis=1)
X_test_plus_y = pd.concat((X_test, y_test), axis=1)
ss = StandardScaler()
ss.fit(X_train)
```

Out[26]:

```
StandardScaler()
```

Created a `OneHotEncoder` for our initial model

In [27]:

```
ohe = OneHotEncoder()
ohe.fit_transform(X_train)
```

```
<17277x26611 sparse matrix of type '<class 'numpy.float64'>'
```

Out[27]: with 172770 stored elements in Compressed Sparse Row format>

Created a training data sat based on the StandardScalar values and our training/test split. We reset the index at the end of each training/test split to avoid any join conflicts/creation of NaNs that would increase number of entries in our dataset.

```
In [28]: # Training dataset
X_train_scaled = pd.DataFrame(ss.transform(X_train))
X_train_scaled.columns = X.columns
y_train.reset_index(drop=True,inplace=True)
X_train_scaled_plus_y = pd.concat((X_train_scaled, y_train), axis=1)
X_train.reset_index(drop=True, inplace=True)
```

Created the test dataset based on our StandardScalar values and our training/test split

```
In [29]: # Testing dataset
X_test_scaled = pd.DataFrame(ss.transform(X_test))
X_test_scaled.columns = X.columns
y_test.reset_index(drop=True,inplace=True)
X_test_scaled_plus_y = pd.concat((X_test_scaled,y_test),axis=1)
X_test.reset_index(drop=True, inplace=True)
```

## First Model - our initial multilinear regression model

Created our first model based on the values in our heatmap that had the highest correlation with price, in this case that was the sqft\_living + bedrooms + floors columns. We will then use this first model to check and see if the predictors have any significant impact on the price of the homes.

```
In [30]: # Keep first model as sqft_living + bedrooms + floors to show that was our first model
# 2nd model should be 'price ~ sqft_living + bedrooms + sqft_per_bedroom + footprint' t
formula = 'price ~ sqft_living + bedrooms + floors'
model = ols(formula, X_train_scaled_plus_y).fit()
model.summary()
```

Out[30]:

OLS Regression Results						
<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.505			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.505			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	5874.			
<b>Date:</b>	Fri, 18 Feb 2022	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	15:26:56	<b>Log-Likelihood:</b>	-2.3986e+05			
<b>No. Observations:</b>	17277	<b>AIC:</b>	4.797e+05			
<b>Df Residuals:</b>	17273	<b>BIC:</b>	4.798e+05			
<b>Df Model:</b>	3					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Intercept</b>	5.409e+05	1970.511	274.521	0.000	5.37e+05	5.45e+05

<b>sqft_living</b>	2.882e+05	2530.414	113.894	0.000	2.83e+05	2.93e+05
<b>bedrooms</b>	-5.411e+04	2402.875	-22.520	0.000	-5.88e+04	-4.94e+04
<b>floors</b>	1412.8026	2110.123	0.670	0.503	-2723.252	5548.857

<b>Omnibus:</b>	11809.269	<b>Durbin-Watson:</b>	1.991
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	435749.399
<b>Skew:</b>	2.803	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	26.956	<b>Cond. No.</b>	2.11

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

After running the model, we see that sqft\_living and bedrooms are a significant factor in determining the price of a home. Since the p-value of floors > 0.05, we can say that floors are not a significant factor in determining the price of a home. This pertains to our business question since the lots of land themselves are expensive, but homes are cheaper to build vertically.

Once again calling the function `sm_metrics` that we defined earlier to find the metrics of our model. We prioritize/focus on RMSE because it provides us with more accuracy. The values in this case make sense since we are dealing with homes that are in the hundreds of thousands range.

```
In [31]: sm_metrics(model, y_test, X_test_scaled)
```

```
Metrics:
Mean Absolute Error: 169584.127
Mean Squared Error: 64475389966.508
Root Mean Squared Error: 253920.046
```

We then calculate the predicted values of our models to calculate the residuals and see if our models were implemented correctly.

```
In [32]: y_hat = model.predict(X_train_scaled)
         y_pred = model.predict(X_test_scaled)
```

We calculate the residuals


```
In [33]: residuals = (y_train - y_hat)
         print(residuals)
```

```
0      134054.190256
1     -177151.427931
2     -39418.666228
3      211972.734232
4     -284322.633742
...
17272  -108179.962734
17273  -18228.613749
```

```
17274    -141713.668295
17275     -61076.700318
17276      34848.286680
Length: 17277, dtype: float64
```

Then we plot on a qqplot to visualize the residuals and see where our error lies. We see that some of our residuals break off towards the left and right side of the plot; however, for the most part, most of our residuals lie along the line.


```
In [34]: qq(X_test_scaled, X_train_scaled, y_test, y_train)
```



Plotted a heteroskedasticity plot for visualization of any recognizable patterns to check if our model was ran correctly.

```
In [35]: plt.figure(figsize = (20,15))
plt.scatter(y_hat, y_train - y_hat, label = 'Training Set')
plt.scatter(y_pred, y_test - y_pred, label = 'Test Set')

plt.axhline(y = 0, color = 'red', label = '0')
plt.xlabel('Predictions')
plt.ylabel('Residuals')
plt.legend()
plt.show()
```



We see after plotting the heteroskedasticity of our plots that a funnel pattern is semi-recognizable. This means that our data is (i don't remember what to say here), so we have to log our data in order to get a normalized plot, and so we will apply that to our second model.

In [36]: `X_train_scaled.corr()`

Out[36]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	yr_built	zipcode	sqft_per
<b>bedrooms</b>	1.000000	0.509423	0.571636	0.026344	0.178529	0.154396	-0.152053	
<b>bathrooms</b>	0.509423	1.000000	0.752564	0.085556	0.503355	0.511066	-0.205060	
<b>sqft_living</b>	0.571636	0.752564	1.000000	0.170260	0.356378	0.320881	-0.201041	
<b>sqft_lot</b>	0.026344	0.085556	0.170260	1.000000	-0.001869	0.053034	-0.128976	
<b>floors</b>	0.178529	0.503355	0.356378	-0.001869	1.000000	0.488497	-0.058795	
<b>yr_built</b>	0.154396	0.511066	0.320881	0.053034	0.488497	1.000000	-0.347291	
<b>zipcode</b>	-0.152053	-0.205060	-0.201041	-0.128976	-0.058795	-0.347291	1.000000	
<b>sqft_per_bedroom</b>	-0.024644	0.517643	0.759121	0.194870	0.293768	0.261316	-0.124074	
<b>footprint</b>	0.030371	0.289180	0.077961	-0.250546	0.553484	0.279235	0.177918	
<b>bedbathratio</b>	0.217513	-0.653090	-0.324053	-0.049564	-0.373462	-0.440836	0.126029	

We then check the VIF to see if we missed anything about multicollinearity between the variables.

```
In [37]: # Check the VIFs after running the model to see if any values are > 10
vif(X_train_scaled)
```

	VIF	features
0	8.133196	bedrooms
1	10.414278	bathrooms
2	15.034290	sqft_living
3	1.142210	sqft_lot
4	1.955622	floors
5	1.739823	yr_built
6	1.276777	zipcode
7	9.195762	sqft_per_bedroom
8	1.764403	footprint
9	6.343995	bedbathratio

Our initial model produced an  $R^2$  value of 0.505. Of our initial predictors, we identified that `sqft_living` and `bedrooms` were significant in accounting for the variation in price while `floors` was **NOT** significant in the variation of price (due to its extremely high P value). In our next model, we are going to identify the predictors that we engineered.

## Second Model - reducing complexity

We engineered this feature so as to reduce the complexity of this model and get rid of discrete categorical variables. We wanted to see if removing these variables and reducing the complexity would increase the efficiency of our model.

- Some of the initial predictors that we had in our initial dataset had multicollinearity with each other so we decided to remove them from our 2nd model.

```
In [38]: formula2 = 'price ~ sqft_living + sqft_per_bedroom + footprint + bedbathratio'
model2 = ols(formula2, X_train_plus_y).fit()
model2.summary()
```

Out[38]:

OLS Regression Results						
<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.502			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.502			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4354.			
<b>Date:</b>	Fri, 18 Feb 2022	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	15:26:57	<b>Log-Likelihood:</b>	-2.3992e+05			
<b>No. Observations:</b>	17277	<b>AIC:</b>	4.798e+05			
<b>Df Residuals:</b>	17272	<b>BIC:</b>	4.799e+05			
<b>Df Model:</b>	4					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Intercept</b>	-1.93e+05	1.26e+04	-15.347	0.000	-2.18e+05	-1.68e+05

<b>sqft_living</b>	240.5972	3.365	71.492	0.000	234.001	247.194
<b>sqft_per_bedroom</b>	248.7397	16.059	15.489	0.000	217.263	280.217
<b>footprint</b>	1.062e+05	7706.949	13.781	0.000	9.11e+04	1.21e+05
<b>bedbathratio</b>	2.561e+04	3731.289	6.865	0.000	1.83e+04	3.29e+04

<b>Omnibus:</b>	12444.927	<b>Durbin-Watson:</b>	1.989
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	551350.015
<b>Skew:</b>	2.973	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	30.029	<b>Cond. No.</b>	1.61e+04

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.61e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In [39]:

```
# Check the VIFs after running the model to see if any values are > 10
vif(X_train_scaled[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedbathratio']])
```

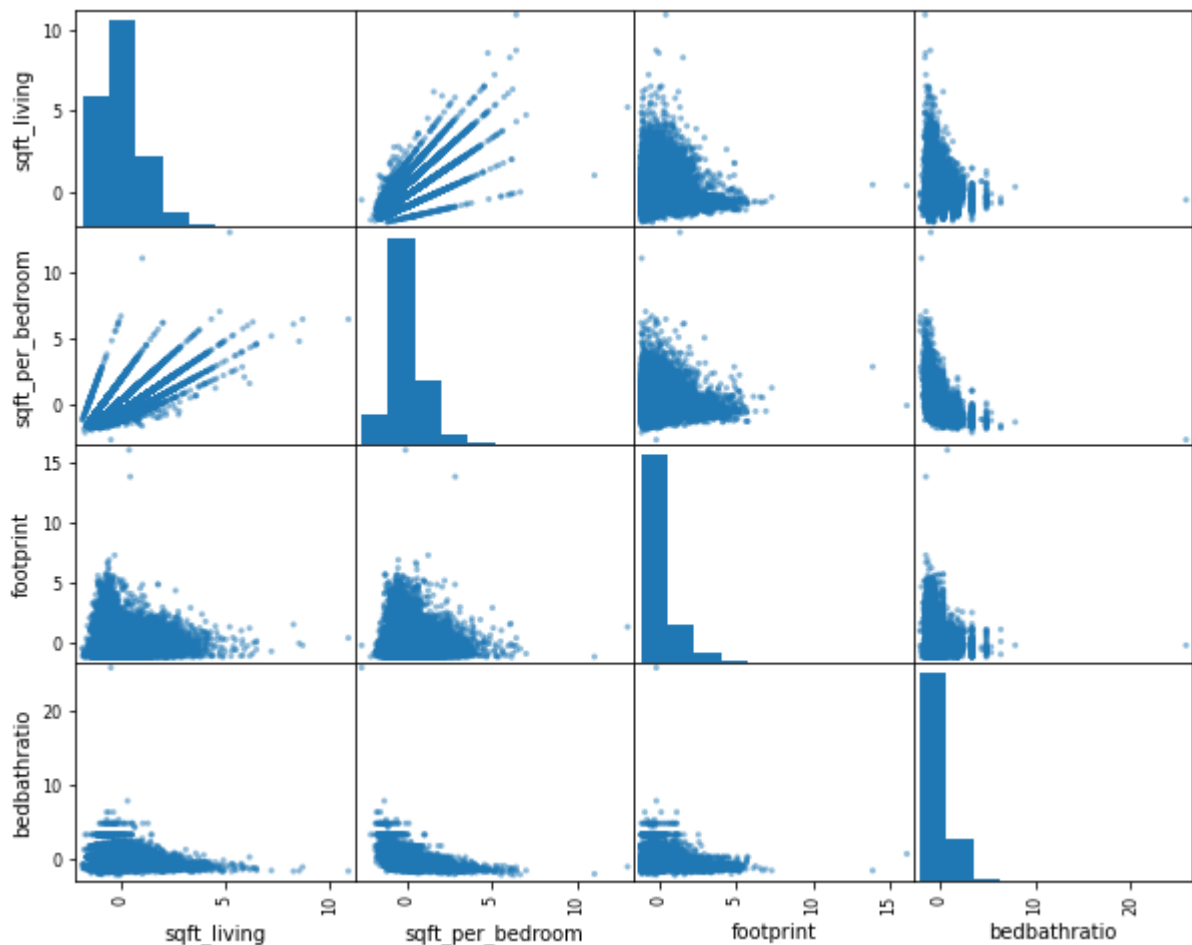
	VIF	features
0	2.429858	sqft_living
1	3.082085	sqft_per_bedroom
2	1.097728	footprint
3	1.570652	bedbathratio

Created a scatter matrix to check for distribution and heteroskedasticity.

In [40]:


```
pd.plotting.scatter_matrix(X_train_scaled[['sqft_living', 'sqft_per_bedroom', 'footprin
```





We now run the second model to see if there were any changes to the  $R^2$  value. We also wanted to test and see if keeping the bedrooms/bathrooms variables would reduce the complexity of our models. Removing the bedrooms variable in our model reduces our  $R^2$  value by 0.015.

```
In [41]: qq(X_test_scaled[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedbathratio']],
           X_train_scaled[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedbathratio']], y_
```




```
In [42]: y_hat2 = model2.predict(X_train[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedba
y_pred2 = model2.predict(X_test[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedba
```

In [43]:

```
plt.figure(figsize = (20,15))
plt.scatter(y_hat2, y_train - y_hat2, label = 'Training Set')
plt.scatter(y_pred2, y_test - y_pred2, label = 'Test Set')

plt.axhline(y = 0, color = 'red', label = '0')
plt.xlabel('Predictions')
plt.ylabel('Residuals')
plt.legend()
plt.show()
```



After we check for heteroskedasticity, we identified a similar "funnel" pattern in our dataset. This led us to believe that the data is not being scaled correctly, so we took the log of the data and plotted the data below to visualize a more normal distribution of our findings.

In [44]:


```
ylog_test2 = np.log(abs(y_pred2))
ylog_train2 = np.log(abs(y_hat2))
```

In [45]:

```
plt.figure(figsize = (20,15))
plt.scatter(ylog_train2, np.log(abs(y_train)) - ylog_train2, label = 'Training Set')
plt.scatter(ylog_test2, np.log(abs(y_test)) - ylog_test2, label = 'Test Set')

plt.axhline(y = 0, color = 'red', label = '0')
plt.xlabel('Predictions')
plt.ylabel('Residuals')
```

```
plt.legend()  
plt.show()
```



```
In [46]: model2.params
```

```
Out[46]: Intercept          -193006.506017  
sqft_living             240.597198  
sqft_per_bedroom        248.739665  
footprint               106212.265597  
bedbathratio           25613.591120  
dtype: float64
```

We run the defined function `sm_metrics` to compare the RMSE values of our training and test dataset, and we identified that the RMSE values are relatively close between the training dataset and the test dataset. The margin of error is still decently high; however, compared to our base model, our new model is more efficient.

```
In [47]: sm_metrics(model2, y_train, X_train)
```

```
Metrics:  
Mean Absolute Error: 170311.199  
Mean Squared Error: 67468439768.192  
Root Mean Squared Error: 259746.876
```

```
In [48]: sm_metrics(model2, y_test, X_test)
```

```
Metrics:
```

Mean Absolute Error: 168308.110  
Mean Squared Error: 64746046225.220  
Root Mean Squared Error: 254452.444

In conclusion, this 3rd model produced an  $R^2$  value of 0.502 and an RMSE of about 260,000. Compared to our second model, this model is more efficient since we utilize less variables and still get about the same  $R^2$  value.

## Third Model - focusing on sqft\_living and sqft\_per\_bedroom

We recognize that sqft\_living is one of the most important factors that we extrapolated from this model and our dataset. We decided to run a new model that focused primarily on the sqft\_living predictors and the sqft\_per\_bedroom predictor that we engineered.

```
In [49]: formula3 = 'price ~ sqft_living + sqft_per_bedroom'
model3 = ols(formula3, X_train_plus_y).fit()
model3.summary()
```

Out[49]:

OLS Regression Results						
<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.496			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.496			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	8510.			
<b>Date:</b>	Fri, 18 Feb 2022	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	15:27:01	<b>Log-Likelihood:</b>	-2.4001e+05			
<b>No. Observations:</b>	17277	<b>AIC:</b>	4.800e+05			
<b>Df Residuals:</b>	17274	<b>BIC:</b>	4.801e+05			
<b>Df Model:</b>	2					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Intercept</b>	-9.454e+04	6031.466	-15.674	0.000	-1.06e+05	-8.27e+04
<b>sqft_living</b>	245.5396	3.335	73.615	0.000	239.002	252.077
<b>sqft_per_bedroom</b>	201.1258	14.132	14.232	0.000	173.425	228.826
<b>Omnibus:</b>	12274.680	<b>Durbin-Watson:</b>	1.990			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	520725.641			
<b>Skew:</b>	2.925	<b>Prob(JB):</b>	0.00			
<b>Kurtosis:</b>	29.251	<b>Cond. No.</b>	7.16e+03			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.16e+03. This might indicate that there are strong multicollinearity or other numerical problems.

After running our model, we get a  $R^2$  value of 0.496. The fact that model produces an  $R^2$  value of 0.496 with only two predictors signifies that sqft has a significant impact on the variation of the price of a home. The coefficients of our predictors in this model are reasonable numbers.


```
In [50]: sm_metrics(model3, y_train, X_train)
```

```
Metrics:  
Mean Absolute Error: 172693.467  
Mean Squared Error: 68246302116.105  
Root Mean Squared Error: 261239.932
```

```
In [51]: sm_metrics(model3, y_test, X_test)
```

```
Metrics:  
Mean Absolute Error: 170805.475  
Mean Squared Error: 65423939839.562  
Root Mean Squared Error: 255781.039
```


```
In [52]: qq(X_test_scaled[['sqft_living', 'sqft_per_bedroom']],  
           X_train_scaled[['sqft_living', 'sqft_per_bedroom']], y_test, y_train)
```



```
In [53]: y_hat3 = model3.predict(X_train[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedba  
y_pred3 = model3.predict(X_test[['sqft_living', 'sqft_per_bedroom', 'footprint', 'bedba
```

```
In [54]: plt.figure(figsize = (20,15))  
plt.scatter(y_hat3, y_train - y_hat3, label = 'Training Set')  
plt.scatter(y_pred3, y_test - y_pred3, label = 'Test Set')  
  
plt.axhline(y = 0, color = 'red', label = '0')  
plt.xlabel('Predictions')  
plt.ylabel('Residuals')
```

```
plt.legend()
plt.show()
```



```
In [55]: ylog_test3 = np.log(abs(y_pred3))
ylog_train3 = np.log(abs(y_hat3))
```

```
In [56]: plt.figure(figsize = (20,15))
plt.scatter(ylog_train3, np.log(abs(y_train)) - ylog_train3, label = 'Training Set')
plt.scatter(ylog_test3, np.log(abs(y_test)) - ylog_test3, label = 'Test Set')

plt.axhline(y = 0, color = 'red', label = '0')
plt.xlabel('Predictions')
plt.ylabel('Residuals')
plt.legend()
plt.show()
```



## Fourth Model - implementing categorical variables

This model implements the categorical variables 'grade' and 'zipcode' that we identified in our datasets. We want to run this model to see if the categorical predictors that we identified from our dataset have a significant influence on the price of a home or not.

```
In [57]: y1 = dfrev2['price']
X1 = dfrev2[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'yr_built', 'zip
X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_s
```

Created a DummyRegressor in similar fashion to the one made for the three models above.

```
In [58]: X1_train_cont = X1_train.drop(['zipcode', 'grade'], axis = 1).copy()
X1_test_cont = X1_test.drop(['zipcode', 'grade'], axis = 1).copy()
ss1 = StandardScaler()
ss1.fit(X1_train_cont)

X1_train_scaled = ss1.transform(X1_train_cont)
X1_test_scaled = ss1.transform(X1_test_cont)

#Dummy Regressor

basemodel1 = DummyRegressor(strategy = 'mean')
basemodel1.fit(X1_train, y1_train)
basemodel1.predict(X1_train)
```

```
array([540946.37500724, 540946.37500724, 540946.37500724, ...,
```

```
Out[58]: 540946.37500724, 540946.37500724, 540946.37500724])
```

Dropped the numerical columns from our dataframe so that we could only focus on and OneHotEncode the categorical values we pulled from our demographics XLS

```
In [59]: X1_cat = X1_train[['zipcode', 'grade']]
ohe1 = OneHotEncoder(sparse = False)
ohe_df = pd.DataFrame(ohe1.fit_transform(X1_cat))
```

```
In [60]: ohe_df.columns = ohe1.get_feature_names(['zipcode', 'grade'])
```

```
In [61]: X1_train_df = pd.DataFrame(X1_train_scaled)
X1_train_df.columns = X1_train_cont.columns.tolist()
```

```
In [62]: X1_train_df
```

```
Out[62]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	yr_built
0	-0.404166	-0.148210	-0.591864	-0.174771	-0.917505	1.123192
1	0.670108	0.832246	0.489562	-0.020598	0.937801	0.884897
2	-0.404166	0.178609	-0.307853	-0.201845	0.937801	0.408307
3	0.670108	0.505427	0.631567	-0.274556	0.937801	-2.076770
4	0.670108	0.832246	2.128085	-0.115790	0.937801	-0.953379
...	...	...	...	...	...	...
17272	-1.478441	0.178609	-0.482629	-0.309172	0.937801	0.306180
17273	1.744382	1.485883	1.319747	-0.208081	0.937801	1.395529
17274	-0.404166	1.485883	1.778533	0.567548	0.937801	0.918939
17275	0.670108	-0.475029	-0.602787	-0.151807	0.010148	-0.102325
17276	-1.478441	-1.455485	-1.028803	-0.161718	-0.917505	-0.783169

17277 rows × 6 columns

```
In [63]: ohe_df.head()
```

```
Out[63]:
```

	zipcode_Northeast	zipcode_Shoreline	zipcode_Southeast	zipcode_Southwest	zipcode_West	grade_10 Very Good
0	0.0	1.0	0.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	1.0	0.0



	zipcode_Northeast	zipcode_Shoreline	zipcode_Southeast	zipcode_Southwest	zipcode_West	grade_10 Very Good
4	0.0	1.0	0.0	0.0	0.0	0.0

Created a dataframe with the new OHE values for our categorical columns from housedf, and we added the price column from our original dataframe to the OHE dataframe to check for correlation between these variables and if they had an effect on home pricing.

```
In [64]: combined_df = ohe_df.join(X1_train_df)
X_train_predictors = combined_df.copy()
combined_df['price'] = y_train
combined_df
```

```
Out[64]:
```

	zipcode_Northeast	zipcode_Shoreline	zipcode_Southeast	zipcode_Southwest	zipcode_West	grade_10 Very Good
0	0.0	1.0	0.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	1.0	0.0
4	0.0	1.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...
17272	0.0	1.0	0.0	0.0	0.0	0.0
17273	1.0	0.0	0.0	0.0	0.0	0.0
17274	1.0	0.0	0.0	0.0	0.0	0.0
17275	0.0	0.0	1.0	0.0	0.0	0.0
17276	0.0	1.0	0.0	0.0	0.0	0.0

17277 rows × 22 columns


```
In [65]: # Before running our model, run correlation to check for multicollinearity.
combined_df.corr()
```

```
Out[65]:
```

	zipcode_Northeast	zipcode_Shoreline	zipcode_Southeast	zipcode_Southwest	zipcode_West
zipcode_Northeast	1.000000	-0.229578	-0.392327	-0.317999	-0.000000
zipcode_Shoreline	-0.229578	1.000000	-0.136507	-0.110646	-0.000000
zipcode_Southeast	-0.392327	-0.136507	1.000000	-0.189083	-0.000000


	zipcode_Northeast	zipcode_Shoreline	zipcode_Southeast	zipcode_Southwest	zipcode_West
zipcode_Southwest	-0.317999	-0.110646	-0.189083	1.000000	-0.199686
zipcode_West	-0.414327	-0.144162	-0.246360	-0.199686	1.000000
grade_10 Very Good	0.164717	-0.055653	-0.077874	-0.067689	-0.067689
grade_11 Excellent	0.117355	-0.037506	-0.054909	-0.041655	-0.041655
grade_12 Luxury	0.053578	-0.007378	-0.028241	-0.016679	-0.016679
grade_13 Mansion	0.007630	0.019168	-0.012191	-0.009881	-0.009881
grade_4 Low	0.001459	-0.010541	0.013731	0.008289	-0.013731
grade_5 Fair	-0.019479	-0.014841	0.021668	0.039150	-0.013731
grade_6 Low Average	-0.162514	0.040858	0.028627	0.214182	-0.013731
grade_7 Average	-0.159235	0.067950	0.100784	0.020687	0.013731
grade_8 Good	0.043572	-0.013586	-0.026789	-0.065107	0.013731
grade_9 Better	0.158288	-0.060654	-0.066623	-0.077702	-0.013731
bedrooms	0.175478	-0.038443	0.010241	-0.093382	-0.013731
bathrooms	0.273140	-0.103396	-0.060813	-0.154872	-0.013731
sqft_living	0.317908	-0.083594	-0.085776	-0.152231	-0.013731
sqft_lot	0.110997	-0.042083	0.048762	-0.029336	-0.013731
floors	0.108523	-0.094373	-0.090932	-0.127626	0.013731
yr_built	0.360141	-0.092777	0.090581	-0.137272	-0.013731
price	0.250924	-0.061293	-0.304331	-0.188512	0.013731

22 rows × 22 columns



In [66]:

```
# Run a heatmap to better visualize the correlation between the variables that we identified
plt.figure(figsize = (25, 20))
sns.heatmap(combined_df.corr().abs(), annot=True);
```



As we can see from the `df_corrs` and the heatmap generated above, there is some correlation between the categorical variables we identified (grade/zipcode/floors) and their effect on price.

After we check the correlation and lack of multicollinearity, we can now run a model on the variables on these variables.

```
In [67]: lr = LinearRegression()
model4 = lr.fit(X_train_predictors, y_train)
model4.score(X_train_predictors, y_train)
```

```
Out[67]: 0.7016031747477318
```

```
In [68]: y_pred_model = lr.predict(X_train_predictors.values)
```

```
In [69]: sk_metrics(y_train, y_pred_model)
```

```
Metrics:
R2: 0.702
Mean Absolute Error: 118807.607
Mean Squared Error: 40430372946.267
Root Mean Squared Error: 201073.054
```

```
In [70]: lr_model_metrics(X1_test, X1_train, y1_test, y1_train, ['grade', 'zipcode']);
```

Train Data

Metrics:

R2: 0.702

Mean Absolute Error: 118783.298

Mean Squared Error: 40429737661.469

Root Mean Squared Error: 201071.474

Test Data

Metrics:

R2: 0.702

Mean Absolute Error: 118921.240

Mean Squared Error: 39508992531.045

Root Mean Squared Error: 198768.691

## Fifth Model - incorporating features from previous model and just the discrete 'grade' feature

```
In [71]: y5 = dfrev2['price']
X5 = dfrev2[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'yr_built', 'grade']]
X5_train, X5_test, y5_train, y5_test = train_test_split(X5, y5, test_size=0.2, random_s
```

```
In [72]: # Define cat_vars
cat_vars = ['grade']

x5_train = pd.DataFrame()
x5_test = pd.DataFrame()
x5_train = X5_train[cat_vars]
x5_train['ID'] = 1
x5_test = X5_test[cat_vars]
x5_test['ID'] = 1

# One hot encoding cat vars
onehot = OneHotEncoder(sparse=False, handle_unknown = 'ignore')
x_train_cat = pd.DataFrame(onehot.fit_transform(x5_train))
x_train_cat.columns = onehot.get_feature_names(x5_train.columns.tolist())
x_test_cat = pd.DataFrame(onehot.transform(x5_test))
x_test_cat.columns = onehot.get_feature_names(x5_train.columns.tolist())

# Drop ID column
x_train_cat.drop(['ID_1'], axis=1, inplace=True)
x_test_cat.drop(['ID_1'], axis=1, inplace=True)

# Resetting indices
x_train_cat.reset_index(drop=True, inplace=True)
x_test_cat.reset_index(drop=True, inplace=True)
X5_train.reset_index(drop=True, inplace=True)
X5_test.reset_index(drop=True, inplace=True)

# Combine dummied cat vars with non-cat
x_train_df = x_train_cat.join(X5_train.drop(cat_vars, axis = 1))
x_test_df = x_test_cat.join(X5_test.drop(cat_vars, axis = 1))
```

<ipython-input-72-0da3b2a3dac0>:7: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
x5_train['ID'] = 1
```

<ipython-input-72-0da3b2a3dac0>:9: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
x5_test['ID'] = 1
```

In [73]:

```
x_train_df
```

Out[73]:

	grade_10 Very Good	grade_11 Excellent	grade_12 Luxury	grade_13 Mansion	grade_4 Low	grade_5 Fair	grade_6 Low Average	grade_7 Average	grade_8 Good	grade_9 Better
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
...	...	...	...	...	...	...	...	...	...	...
17272	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0
17273	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1
17274	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0
17275	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0
17276	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0

17277 rows × 16 columns



After we ran the model, we checked the score value of the training set vs the test set. We also doublecheck the metrics to see if there are any differences we failed to identify.

In [74]:

```
lr = LinearRegression()  
model5 = lr.fit(x_train_df, y5_train)  
model5.score(x_train_df, y5_train)
```

Out[74]:

```
0.644614517676586
```

In [75]:

```
model5.score(x_test_df, y5_test)
```

Out[75]:

```
0.6448091772630099
```

```
In [76]: sk_metrics(y5_train, model5.predict(x_train_df))
```

```
Metrics:  
R2: 0.645  
Mean Absolute Error: 140500.728  
Mean Squared Error: 48151878217.462  
Root Mean Squared Error: 219435.362
```

```
In [77]: sk_metrics(y5_test, model5.predict(x_test_df))
```

```
Metrics:  
R2: 0.645  
Mean Absolute Error: 140062.150  
Mean Squared Error: 47165636238.606  
Root Mean Squared Error: 217176.509
```

After running this model, we see that the  $R^2$  value is 0.645. This is roughly a 15% increase from the model that we ran without the categorical variables. We recognize that the model is a good model since the RMSE values are not extraneously high. Thus we can conclude that this model is more efficient than


```
In [78]: print(x_test_df.columns)  
print(model5.coef_)
```

```
Index(['grade_10 Very Good', 'grade_11 Excellent', 'grade_12 Luxury',  
      'grade_13 Mansion', 'grade_4 Low', 'grade_5 Fair',  
      'grade_6 Low Average', 'grade_7 Average', 'grade_8 Good',  
      'grade_9 Better', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',  
      'floors', 'yr_built'],  
      dtype='object')  
[-2.69101389e+04  2.34921054e+05  8.03776296e+05  1.81228481e+06  
 -6.13175604e+05 -5.98404401e+05 -5.50234177e+05 -4.63710683e+05  
 -3.68105276e+05 -2.30441881e+05 -3.39039328e+04  5.81524173e+04  
  1.54814698e+02 -2.61255318e-01  2.67150446e+04 -3.85123851e+03]
```

```
In [79]: plt.figure(figsize = (20,15))  
plt.scatter(model5.predict(x_train_df), y5_train - model5.predict(x_train_df), label =  
plt.scatter(model5.predict(x_test_df), y5_test - model5.predict(x_test_df), label = 'Te  
  
plt.axhline(y = 0, color = 'red', label = '0')  
plt.xlabel('Predictions')  
plt.ylabel('Residuals')  
plt.legend()  
plt.show()
```



```
In [80]: qq(x_test_df, x_train_df, y5_test, y5_train)
```



After we ran our model, we ran a qqplot and heteroskedasticity check to see if there were any huge errors in our modeling. We recognize that there is a pattern in most of our data that is likely attributed to the lack of normal distribution between of our data(it is data from only one year after all.)

## Sixth Model - Modeling with just sqft\_living, zipcode, and grade of homes

After we ran the previous model, we wanted to create our last and final model to see the significance of just the three most important predictors that we identified over the course of our analysis and modeling: 'sqft\_living', 'zipcode', and 'grade'.

```
In [82]: y6 = dfrev2['price']
X6 = dfrev2[['sqft_living', 'zipcode', 'grade']]
x6_train, x6_test, y6_train, y6_test = train_test_split(X6, y6, test_size=0.2, random_s
```

```
In [83]: # Onehot encode

# Define cat_vars
cat_vars = ['grade', 'zipcode']

# One hot encoding cat vars
onehot = OneHotEncoder(sparse=False, handle_unknown = 'ignore')
x6_train_cat = pd.DataFrame(onehot.fit_transform(x6_train[cat_vars]))
x6_train_cat.columns = onehot.get_feature_names(cat_vars)
x6_test_cat = pd.DataFrame(onehot.transform(x6_test[cat_vars]))
x6_test_cat.columns = onehot.get_feature_names(cat_vars)

# Resetting indices
x6_train_cat.reset_index(drop=True, inplace=True)
x6_test_cat.reset_index(drop=True, inplace=True)
x6_train.reset_index(drop=True, inplace=True)
x6_test.reset_index(drop=True, inplace=True)

# Combine dummied cat vars with non-cat
x6_train_df = x6_train_cat.join(x6_train.drop(cat_vars, axis = 1))
x6_test_df = x6_test_cat.join(x6_test.drop(cat_vars, axis = 1))
```

```
In [84]: # Run regression

lr = LinearRegression()
model6 = lr.fit(x6_train_df, y6_train)
model6.score(x6_train_df, y6_train)
```

```
Out[84]: 0.6792846894781521
```

```
In [85]: model6.score(x6_test_df, y6_test)
```

```
Out[85]: 0.6803566979982765
```

As you can see, this final model provided us with an  $R^2$  value of 0.680, which is significantly better than our initial model that we produced with just the continuous data. After checking the  $R^2$  value, we need perform a qqplot and heteroskedasticity check to identify normal distribution within our modeling.

```
In [86]: vif(x6_train_df)
```



```
luence.py:193: RuntimeWarning: divide by zero encountered in double_scalars
  vif = 1. / (1. - r_squared_i)
      VIF      features
0      inf  grade_10 Very Good
1      inf  grade_11 Excellent
2      inf   grade_12 Luxury
3      inf  grade_13 Mansion
4      inf   grade_4 Low
5      inf   grade_5 Fair
6      inf grade_6 Low Average
7      inf   grade_7 Average
8      inf   grade_8 Good
9      inf   grade_9 Better
10     inf  zipcode_Northeast
11     inf  zipcode_Shoreline
12     inf  zipcode_Southeast
13     inf  zipcode_Southwest
14     inf   zipcode_West
15  2.583583      sqft_living
```

```
In [87]: sk_metrics(y6_train, model6.predict(x6_train_df))
```

```
Metrics:
R2: 0.679
Mean Absolute Error: 123400.864
Mean Squared Error: 43454348426.844
Root Mean Squared Error: 208457.066
```

```
In [88]: sk_metrics(y6_test, model6.predict(x6_test_df))
```

```
Metrics:
R2: 0.680
Mean Absolute Error: 122776.053
Mean Squared Error: 42445296283.693
Root Mean Squared Error: 206022.563
```

Once again we checked the metrics and we can see that the RMSE of this last model is significantly lower than the RMSE of our baseline model. Since these values are relatively low, we can conclude that our models are not overfit on the dataset and that the **predictors 'sqft\_living', 'zipcode', and 'grade' account for ~68% of the variation in the price of a home.**

```
In [89]: print(x6_test_df.columns)
print(model6.coef_)
```

```
Index(['grade_10 Very Good', 'grade_11 Excellent', 'grade_12 Luxury',
      'grade_13 Mansion', 'grade_4 Low', 'grade_5 Fair',
      'grade_6 Low Average', 'grade_7 Average', 'grade_8 Good',
      'grade_9 Better', 'zipcode_Northeast', 'zipcode_Shoreline',
      'zipcode_Southeast', 'zipcode_Southwest', 'zipcode_West',
      'sqft_living'],
      dtype='object')
[-1.21259720e+05  1.27964757e+05  6.96803430e+05  1.74215418e+06
 -4.08425392e+05 -4.41443189e+05 -4.28565748e+05 -4.44161409e+05
 -4.10033064e+05 -3.13033841e+05  2.06300341e+04  1.09723229e+04
 -1.59947475e+05 -6.95254609e+04  1.97870579e+05  1.78539231e+02]
```


We once again plot the heteroskedasticity and qqplot to check for normal distribution within our

predictors and dataframe.

In [90]:


```
plt.figure(figsize = (20,15))
plt.scatter(model5.predict(x_train_df), y5_train - model5.predict(x_train_df), label = 'Training Set')
plt.scatter(model5.predict(x_test_df), y5_test - model5.predict(x_test_df), label = 'Test Set')

plt.axhline(y = 0, color = 'red', label = '0')
plt.xlabel('Predictions')
plt.ylabel('Residuals')
plt.legend()
plt.show()
```



In [91]:

```
qq(x6_test_df, x6_train_df, y6_test, y6_train)
```



## Visualizations

We decided to extrapolate the data that we produced from our modeling and created bargraphs for easier visualization of our main points about the sqft\_living, zipcode(region), grade(condition) and their significance in home pricing.

```
In [92]: # Binning living space by house size as per National Association of Home Builder (NAHB)

living_price = housedf[['sqft_living', 'price']]

bins = [1., 2000., 3000., 4000., 14000]
bin_labels = ['Small House', 'Medium House', 'Large House', 'Bigger Houses']

living_price['bins'] = pd.cut(x=living_price['sqft_living'], bins=bins, labels=bin_labels)
```

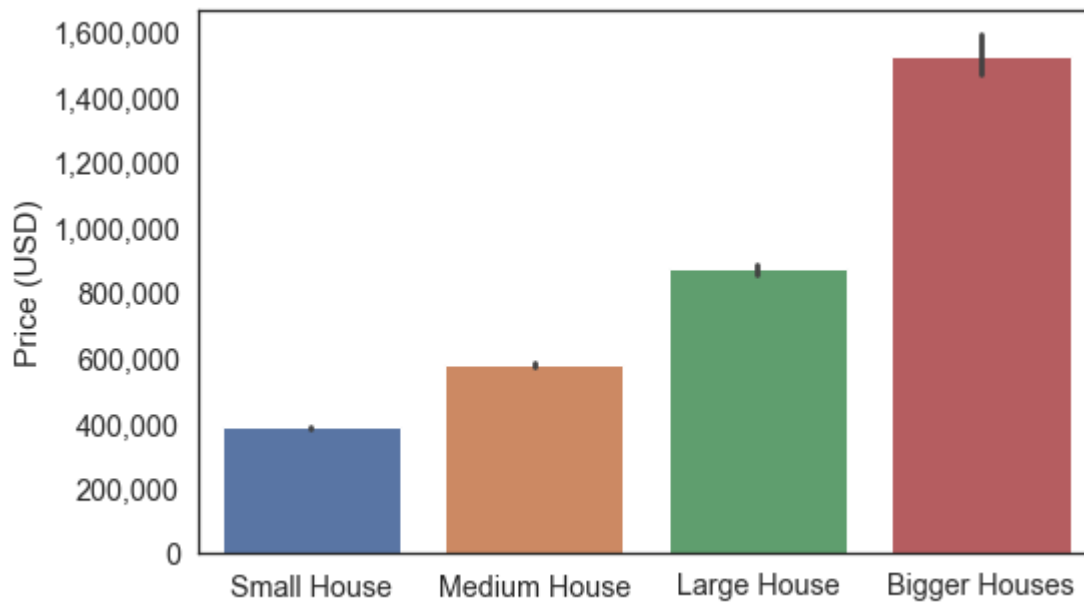
<ipython-input-92-44b9be3c1b63>:8: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
living_price['bins'] = pd.cut(x=living_price['sqft_living'], bins=bins, labels=bin_labels)
```

```
In [93]: # Plot living area bins versus price
plt.figure(figsize=(8, 5))
sns.set(font_scale=1.3, style="white")
v = sns.barplot(x='bins', y='price', data=living_price)
v.yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))
v.set_xlabel("")
v.set_ylabel("Price (USD)")
```


```
Out[93]: Text(0, 0.5, 'Price (USD)')
```



```
In [94]: # Plot price by region
region_price = dfrev2[['zipcode', 'price']]

plt.figure(figsize=(8, 5))
sns.set(font_scale=1.3, style="white")
v = sns.barplot(x='zipcode', y='price', data=region_price)
v.yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))
v.set_xlabel("Court District")
v.set_ylabel("Price (USD)")
```

```
Out[94]: Text(0, 0.5, 'Price (USD)')
```



```
In [95]: # Plot price by grade

# Bin grades
grade_price = housedf[['price', 'grade']]
```

```

grade_def = {'4 Low': 'Low', '3 Poor': 'Low', '6 Low Average': 'Low', '5 Fair' : 'Low'

grade_price['grade'] = grade_price['grade'].map(grade_def)

grade_price['grade'].value_counts().sum()

```

<ipython-input-95-8c42c9c7613c>:10: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

grade_price['grade'] = grade_price['grade'].map(grade_def)
21597

```

Out[95]:

In [96]:

```


# Plot price by grade

grade_order= ['Low', 'Medium', 'High', 'Outstanding']

plt.figure(figsize=(8, 5))
sns.set(font_scale=1.3, style="white")
v = sns.barplot(x='grade', y='price', data=grade_price, order=grade_order)
v.yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))
v.set_xlabel("Grade")
v.set_ylabel("Price (USD)")

```

Out[96]: Text(0, 0.5, 'Price (USD)')



## Conclusion

Our business stakeholder is Howard S Wright construction company, and we identified **three important factors** that they should while building a home:

- **Square Footage**
- **Region to build in(zipcode)**
- **Quality of the home**

These factors are the most significant predictors in maximizing their profit on building a home.

We ran a multilinear regression models to identify that these predictors/factor account for **68% of the variation in the price of home** in King County.